

Jason Teixeira

RAG & Evaluation Engineer | LLM Quality, Eval Harnesses, Python, TypeScript
Orlando, FL | Remote | (908) 937-6568 | nexuraltrading@gmail.com

sageideas.dev | nexural.io | github.com/JasonTeixeira | linkedin.com/in/jason-teixeira

SUMMARY

AI evaluation-focused engineer with production experience building metric-grounded GenAI / LLM workflows, RAG-style source grounding, response validation, API systems, automation gates, and test-heavy applications. Strong in Python, TypeScript, FastAPI, Next.js, Postgres, pytest, Playwright, CI/CD, and AI quality patterns including eval harnesses, golden cases, prompt engineering, prompt regression, hallucination checks, source grounding, and LLM output validation.

TECHNICAL SKILLS

AI Evaluation: GenAI, eval harnesses, golden datasets, LLM-as-judge patterns, response validation, hallucination detection, groundedness, prompt engineering, prompt regression, confidence scoring, semantic checks

RAG: retrieval augmented generation, embeddings, vector search, chunking, reranking, source attribution, knowledge workflows, pgvector-ready Postgres architectures

Languages: Python, TypeScript, JavaScript, C#/.NET, SQL

Backend: FastAPI, Node.js, ASP.NET Core, REST APIs, WebSocket, API contracts

Data/ML: PostgreSQL, Supabase, SQLite, Redis, pandas, scikit-learn, model training/evaluation, machine learning basics, statistical analysis, data pipelines

Testing: pytest, Playwright, API testing, E2E testing, contract tests, regression testing, coverage gates

DevOps: Vercel, AWS, Docker, Kubernetes, GitHub Actions, CI/CD

PROFESSIONAL EXPERIENCE

AI Evaluation Engineer / Founder & Lead Developer | Nexural

2024-Present | Remote

- Built AI response validation that cross-references model claims against actual strategy metrics and confidence scores before exposing AI analysis to users.
- Built Nexural Research with 3 AI providers, full 71-metric context injection, multi-turn strategy conversations, validated AI analysis, 58 FastAPI endpoints, 604 tests, and 93% coverage.
- Designed deterministic metric pipelines and model-facing context previews so users can inspect exactly what the AI receives before asking questions.
- Built Nexural Automation, an MCP automation server with 8 stable tools, golden contract fixtures, schema validation, contract tests, secret scanning, CI quality gates, and repeatable tool outputs.
- Built quality gates across dependency audits, Docker builds, Trivy scans, schema validation, test suites, and release checks to prevent regressions in automation workflows.
- Built Vercel-deployed Next.js/Supabase applications with Playwright E2E tests, RLS isolation checks, API routes, serverless functions, audit logs, and production deployment metadata.

Senior Analyst & Full-Stack Developer | HighStrike

2021-2026 | Remote

- Built fintech dashboards and backend APIs for market data, research workflows, trade ingestion, billing, and subscription automation.
- Authored 200+ pages of financial analysis, technical documentation, research reports, and user-facing explanations requiring precise handling of complex data.
- Collaborated with remote product, engineering, and executive stakeholders using Agile workflows, Git PRs, code reviews, Jira, Slack, and async communication.

SELECTED PROJECTS

Nexural Research: Multi-provider AI analysis, response validation, 604 tests, 93% coverage, FastAPI, Docker/Kubernetes, 71+ metrics.

Nexural Automation: MCP tools, golden fixtures, contract tests, schema validation, agent-callable workflows, secret scanning, CI gates.

Sage Ideas Studio: Production Next.js/Supabase/Stripe app with 105 routes, 57 APIs, RLS, audit logs, Playwright, Vercel deployments.

EDUCATION & CERTIFICATIONS

B.S. Computer Science | Full Sail University

B.S. Finance | Kean University

ISTQB CTFL | ISTQB Test Automation Engineer | ISTQB CT-AI | AWS Cloud Practitioner | AWS Solutions Architect | AWS ML Foundations